

A Bug Hunter's Perspective on Linux Drivers



A walk through the land of I/O Control

Agenda

I. Introduction

II. Device Drivers

- I. I/O Control

- II. Talking to a driver

III. Bug Hunting

- I. Discovery & Debugging

- II. Tooling

IV. Conclusion

#whoami

- Jeremy Brown
 - Independent researcher / consultant
 - Formerly of Microsoft
 - Windows/Phone/Xbox Security
 - Malware Protection Center
 - Also, Tenable
 - Nessus
 - RE patches

What I won't teach you

- Comprehensive driver fundamentals
 - You need lots of time and a few different books
- How to become a driver security expert
 - Only a few of them around– djrbliss, spender, j00ru, etc
- Probably how to earn a living from these bugs
 - It's OK to do research you enjoy, not just for \$\$\$

What I hope to teach you

- What you “need to know” to get started
 - Determining attack surface
 - Finding and reviewing IOCTL handlers
 - Some primitive fuzzing techniques
- Knowledge applicable to *nix device drivers
 - Some of which translates to other OSes
- Take you from knowing nothing to something

Agenda

I. Introduction

II. Device Drivers

I. I/O Control

II. Talking to a driver

III. Bug Hunting

I. Discovery & Debugging

IV. Conclusion

Device drivers

- Plug-in for the OS
 - Code that simply talks to a device
 - Commonly physical or virtual devices
 - Lives in kernel or user space
 - Natively or through abstraction layers
 - eg. wrappers or libraries

Frameworks

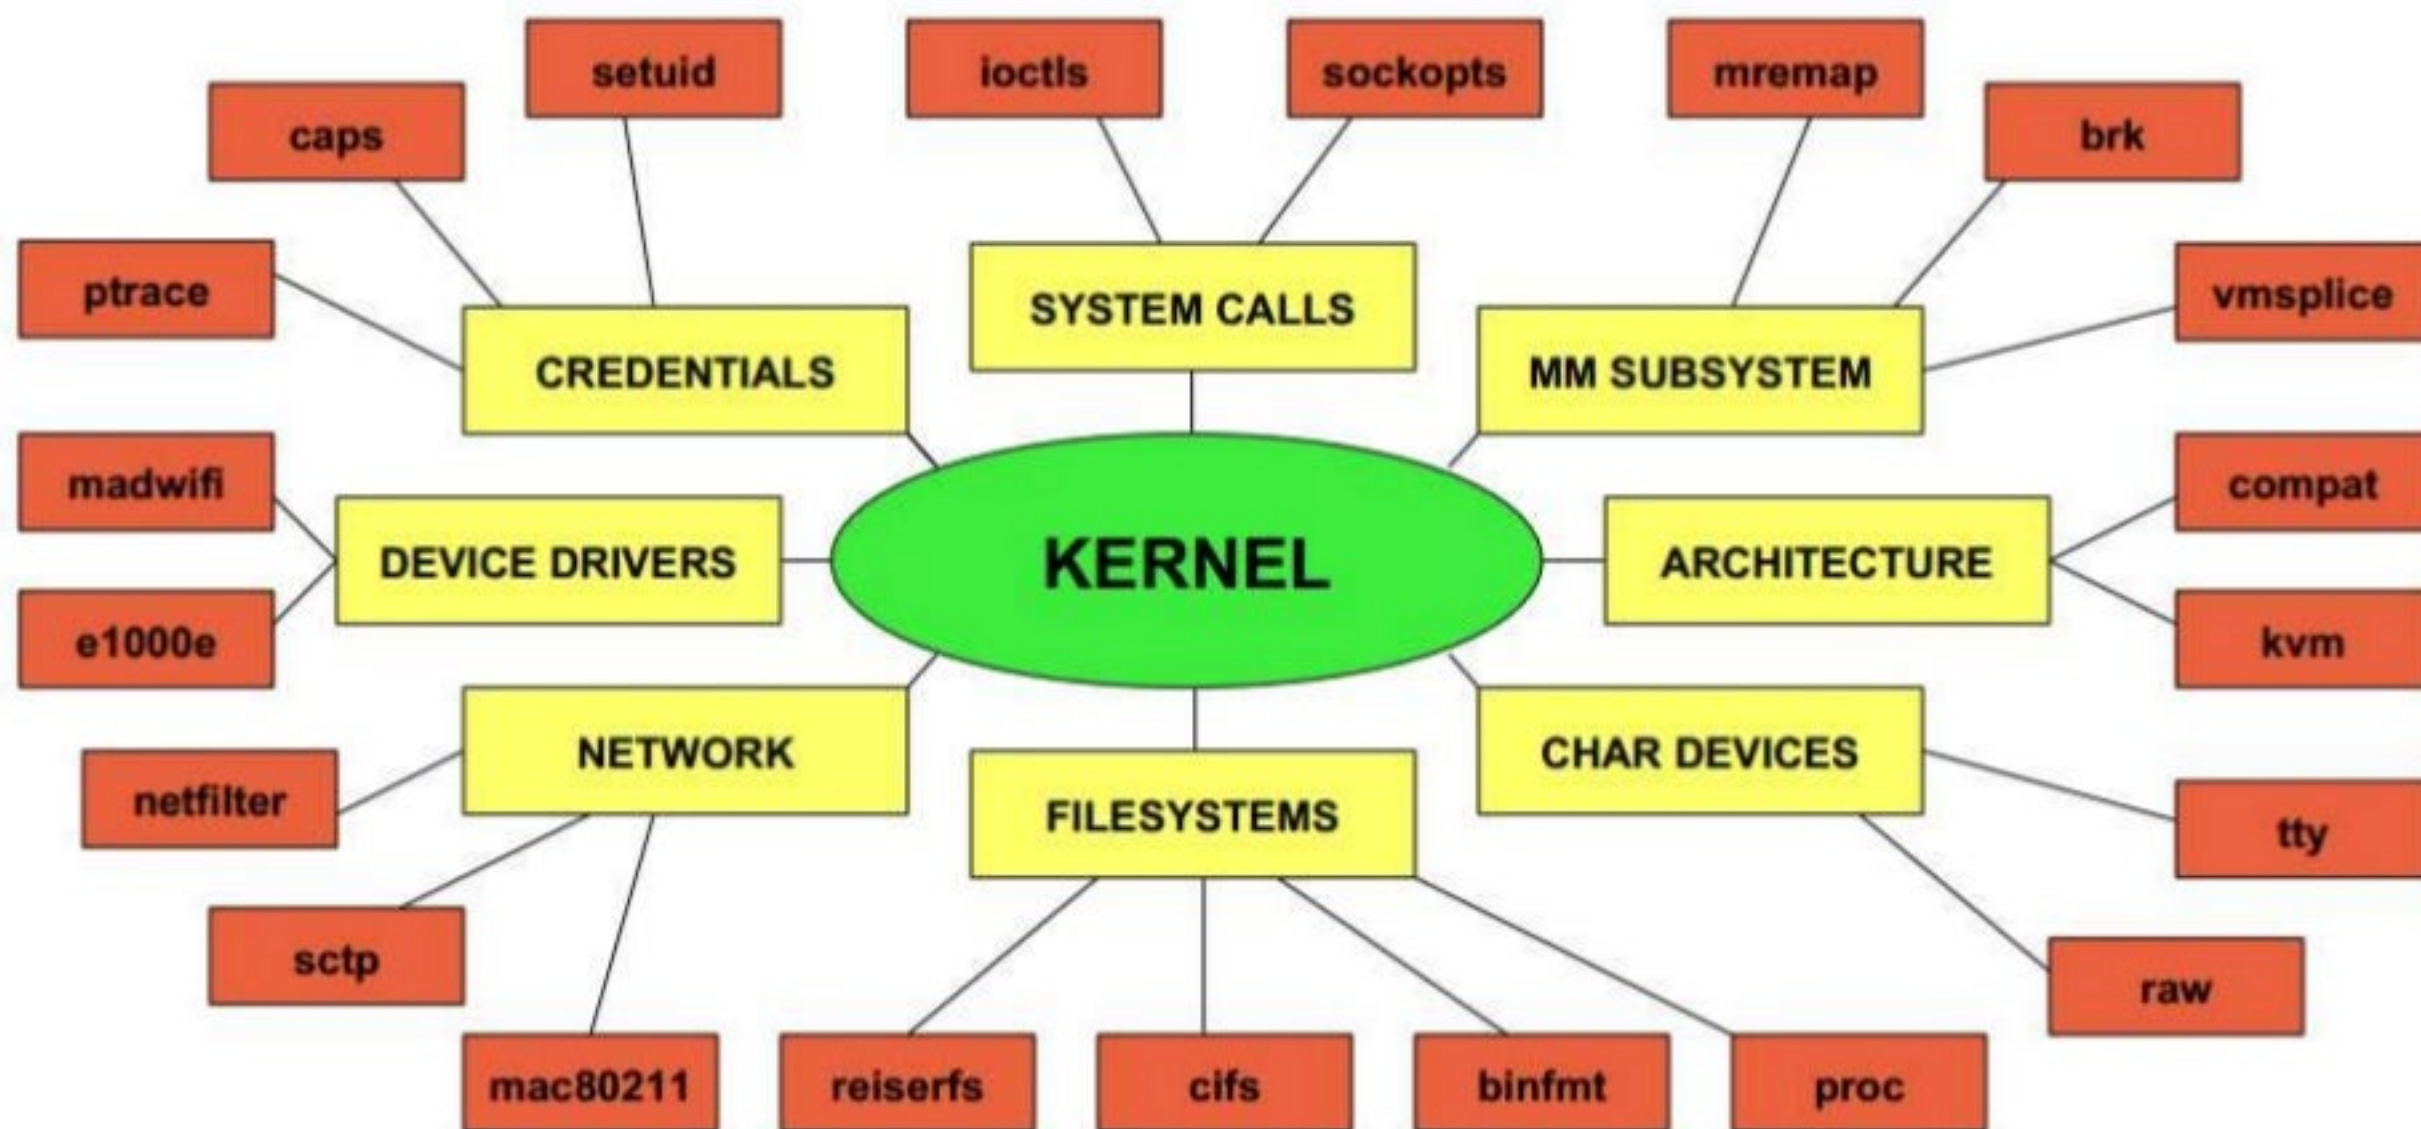
- Traditional Drivers
 - Live in kernel space, eg. HID.ko
 - Cons are kernel panics, debugging & overhead
- Userspace Drivers
 - Live in user space, eg. native frameworks like UIO
 - Interesting con – it can't register interrupt handler
 - Must either poll or use some other driver (kernel)
 - Pro is that network libraries are mature

Where are these drivers?

- **lsmod** shows you which ones are loaded
 - Parses /proc/modules

```
Module          Size  Used by
rfcomm          73557  2
fuse            91453  3
xt_CHECKSUM     12549  1
ipt_MASQUERADE  12678  3
nf_nat_masquerade_ipv4 13203  1 ipt_MASQUERADE
tun             27106  1
nf_conntrack_netbios_ns 12665  0
nf_conntrack_broadcast 12527  1 nf_conntrack_netbios_ns
ip6t_rpfilter   12546  1
ip6t_REJECT     12625  2
nf_reject_ipv6  13301  1 ip6t_REJECT
xt_conntrack    12760  32
ebtable_nat     12807  0
ebtable_broute  12731  0
bridge         108443  1 ebtable_broute
stp             12868  1 bridge
llc             13941  2 stp,bridge
ebtable_filter  12827  0
ebtables        30758  3 ebtable_broute,ebtable_nat,ebtable_filter
```

Linux Kernel Attack Surface



Credit: Jon Oberheide, SOURCE Boston 2010, "Linux Kernel Exploitation"
Reference: <https://jon.oberheide.org/files/source10-linuxkernel-jonoberheide.pdf>

I/O Control

- Common way to operate a device other than simply read/write
 - “We need a way to format floppies and war dial!”
- Implemented by a *special* type of system call
 - Most *nix use **ioctl()**
 - On Windows it's **DeviceIoControl()**

I/O Control

- What do kernel devs say about I/O Control?
 - “The `ioctl()` system call has long been out of favor among the kernel developers, who see it as a **completely uncontrolled entry point into the kernel**[...]”

I/O Control

- Continued
 - “Given the vast number of applications which expect `ioctl()` to be present, however, **it will not go away anytime soon**”



A worthy target

- Drivers provide kernel entry points
 - Find bugs in drivers, execute codes in kernel mode
- Various Impacts
 - Privilege escalation
 - Info Leaks
 - Sandbox escapes

Previous Work

Search for:

Search

Hits per page: 10 ▼

Language: english

[Syntax help](#)

1-10 of exactly 26 matches found in 0.005398 seconds; search the [wiki](#) or the [packages](#)

[Debian -- Security Information -- DSA-2240-1 linux-2.6](#)

Debian Security Advisory DSA-2240-1 linux-2.6 -- privilege escalation/denial of service/information leak Date Reported: 24 May 2011 A 1078, CVE-2011-1079, CVE-2011-1080, CVE-2011-1090, CVE-2011-1160, CVE-2011-1163, CVE-2011-1170, CVE-2011-1171, CVE-100% relevant, matching: *ioctl*

Debian -- Security Information -- DSA-2769-1 kfreebsd-9

Debian Security Advisory DSA-2769-1 kfreebsd-9 -- privilege escalation/denial of service Date Reported: 08 Oct 2013 Affected Package: FreeBSD kernel that may lead to a denial of service or privilege escalation. The Common Vulnerabilities and Exposures project identifies 94% relevant, matching: *ioctf*

Debian -- Security Information -- DSA-2389-1 linux-2.6

Debian Security Advisory DSA-2389-1 linux-2.6 -- privilege escalation/denial of service/information leak Date Reported: 15 Jan 2012 A:
4110, CVE-2011-4127, CVE-2011-4611, CVE-2011-4622, CVE-2011-4914. More information: Several vulnerabilities have been discov
93% relevant, matching: *ioctl*

Debian -- Security Information -- DSA-2126-1 linux-2.6

Debian Security Advisory DSA-2126-1 linux-2.6 -- privilege escalation/denial of service/information leak Date Reported: 26 Nov 2010 A 3432, CVE-2010-3437, CVE-2010-3442, CVE-2010-3448, CVE-2010-3477, CVE-2010-3705, CVE-2010-3848, CVE-2010-3849, CVE-83% relevant, matching: *ioctf*

Previous Work

Date ▼	D	A	V	Title
2014-05-26	📌	-	🕒	Linux kernel 3.14-rc1 <= 3.15-rc4 - Raw Mode PTY Local Echo Race Condition (x64) Local...
2014-02-11	📌	-	🕒	Linux Kernel - Local Root Exploit (ARM)
2011-09-05	📌	-	✅	Linux Kernel < 2.6.36.2 - Eiconet Privilege Escalation Exploit
2011-09-01	📌	-	🕒	Linux Kernel 'perf_count_sw_cpu_clock' event Denial of Service
2011-03-14	📌	-	🕒	Linux <= 2.6.37-rc1 serial_core TIOCGICOUNT Leak Exploit
2011-01-10	📌	-	🕒	Linux Kernel - Solaris < 5.10 138888-01 - Local Root Exploit
2011-01-08	📌	-	🕒	Linux Kernel < 2.6.34 - CAP_SYS_ADMIN x86 & x64 - Local Privilege Escalation Exploit (2)
2011-01-05	📌	-	✅	Linux Kernel 2.6.34 - CAP_SYS_ADMIN x86 - Local Privilege Escalation Exploit
2010-12-07	📌	-	✅	Linux Kernel <= 2.6.37 - Local Privilege Escalation
2010-10-28	📌	-	✅	Linux Kernel - VIDIOCSMICROCODE IOCTL Local Memory Overwrite Vulnerability
2010-09-29	📌	-	✅	Linux Kernel < 2.6.36-rc6 pktcdvd Kernel Memory Disclosure
2010-08-27	📌	-	✅	Linux Kernel < 2.6.36-rc1 CAN BCM - Privilege Escalation Exploit

Reference: <https://www.exploit-db.com/search/?action=search&description=linux&text=ioctl>

Android Privilege Escalation

- Recent blog detailing CVE-2014-4323

The "mdp" driver is extremely complex, supporting a wide range of commands; from IOCTLs, to memory mapping the device, etc.

```
static struct fb_ops mdss_fb_ops = {
    .owner = THIS_MODULE,
    .fb_open2 = mdss_fb_open,
    .fb_release2 = mdss_fb_release,
    .fb_check_var = mdss_fb_check_var, /* vinfo check */
    .fb_set_par = mdss_fb_set_par, /* set the video mode */
    .fb_blank = mdss_fb_blank, /* blank display */
    .fb_pan_display = mdss_fb_pan_display, /* pan display */
    .fb_ioctl = mdss_fb_ioctl, /* perform fb specific ioctl */
    .fb_mmap = mdss_fb_fbmem_ion_mmap,
};
```

This means we need a good strategy for mapping out the **weak spots within the driver**. Skimming over the code, going by the sheer amount of **IOCTL commands supported (at least twenty different commands)**, it seems as though looking at the IOCTL commands in depth might be a lucrative venture.

Funnily, though, there was no need to go too deeply, since the second IOCTL command turned out to be vulnerable :)

Android Privilege Escalation

- Vulnerabilities in the Samsung S4

One year ago, we found some security issues in the Samsung S4 (GT-I9500) version I9500XXUEMK8. After several emails with the Samsung security team, these issues are still unpatched. This blog post *details* these issues and provides a potential patch. The affected driver is the `samsung_extdisp` and CVEs assigned are:


- 1 Kernel memory disclosure (CVE-2015-1800)
- 4 Kernel memory corruption (CVE-2015-1801)

- **memcpy()** doesn't check user/kernel access!

```
case FBIOGET_FSCREENINFO:
-     ret = memcpy(argp, &fb->fix, sizeof(fb->fix)) ? 0 : -EFAULT;
+     ret = copy_to_user(argp, &fb->fix, sizeof(fb->fix)) ? -EFAULT : 0;
    break;

case FBIOGET_VSCREENINFO:
-     ret = memcpy(argp, &fb->var, sizeof(fb->var)) ? 0 : -EFAULT;
+     ret = copy_to_user(argp, &fb->var, sizeof(fb->var)) ? -EFAULT : 0;
    break;
```


Time to Learn



Device types

- Char
 - Character Device
 - Think of it as a “stream of bytes”
 - Mapped to file system nodes, eg. `/dev/XXXXX`
- Char device vs regular file
 - Sequential and arbitrary access, respectively



Device types

- Block
 - Usually hosts a filesystem
 - Also accessible by filesystem nodes
 - But through a different interface than char devices



Device types

- Network
 - Device that exchanges data with hosts
 - Eg. Loopback, eth0
 - No traditional read/write
 - Special packet transmission functions are called



Talking to a driver



Example driver.c

```
#ifndef QUERY_IOCTL_H
#define QUERY_IOCTL_H
#include <linux/ioctl.h>

typedef struct
{
    int status, dignity, ego;
} query_arg_t;

#define QUERY_GET_VARIABLES _IOR('q', 1, query_arg_t *)
#define QUERY_CLR_VARIABLES _IO('q', 2)
#define QUERY_SET_VARIABLES _IOW('q', 3, query_arg_t *)

#endif
```

Example driver.c

```
switch (cmd)
{
    case QUERY_GET_VARIABLES:
        q.status = status;
        q.dignity = dignity;
        q.ego = ego;
        if (copy_to_user((query_arg_t *)arg, &q, sizeof(query_arg_t)))
        {
            return -EACCES;
        }
        break;
    case QUERY_CLR_VARIABLES:
        status = 0;
        dignity = 0;
        ego = 0;
        break;
    case QUERY_SET_VARIABLES:
        if (copy_from_user(&q, (query_arg_t *)arg, sizeof(query_arg_t)))
        {
            return -EACCES;
        }
        status = q.status;
        dignity = q.dignity;
        ego = q.ego;
        break;
    default:
        return -EINVAL;
}
```

copy_to_user()

- Does some overflow checks and then calls
 - `_copy_to_user()`

```
unsigned long _copy_to_user(void __user *to, const void *from, unsigned n)
{
    if (access_ok(VERIFY_WRITE, to, n))
        n = _copy_to_user(to, from, n);
    return n;
}
```

- **access_ok** ensures userspace pointer is valid

copy_from_user()

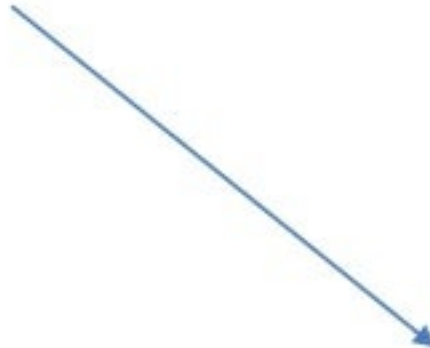
- Calls **_copy_from_user()**

```
unsigned long _copy_from_user(void *to, const void __user *from, unsigned n)
{
    if (access_ok(VERIFY_READ, from, n))
        n = __copy_from_user(to, from, n);
    else
        memset(to, 0, n);
    return n;
}
```

- Same idea, but checks what the user specified for **from** instead of **to**
- Also **VERIFY_READ** instead of **VERIFY_WRITE**

Example client.c

```
char *file_name = "/dev/query";  
int fd;  
enum  
{  
    e_get,  
    e_clr,  
    e_set  
} option;
```



```
fd = open(file_name, O_RDWR);  
if (fd == -1)  
{  
    perror("query_apps open");  
    return 2;  
}  
  
switch (option)  
{  
    case e_get:  
        get_vars(fd);  
        break;  
    case e_clr:  
        clr_vars(fd);  
        break;
```


Example client.c

```
void get_vars(int fd)
{
    query_arg_t q;

    if (ioctl(fd, QUERY_GET_VARIABLES, &q) == -1)
    {
        perror("query_apps ioctl get");
    }
    else
    {
        printf("Status : %d\n", q.status);
        printf("Dignity: %d\n", q.dignity);
        printf("Ego      : %d\n", q.ego);
    }
}

void clr_vars(int fd)
{
    if (ioctl(fd, QUERY_CLR_VARIABLES) == -1)
    {
```

ioctl()

```
#include <sys/ioctl.h>
```

```
int ioctl(int fd, unsigned long request, ...);
```

- 1st arg
 - an opened file descriptor for the device
- 2nd arg
 - control code, specific to each operation
- 3rd arg (optional)
 - Buffer to send or receive data

IOCTL Macros

- IO
 - No data transfer
- IOR
 - Read by kernel, write to user
- IOW
 - Read from user, write by kernel
- IORW
 - Bi-directional data transfer

IOCTL Macros

- Four bitfields
 - Type (magic number)
 - Number
 - Direction
 - Size

```
#define _IO(type,nr)          _IOC(_IOC_NONE,(type),(nr),0)
#define _IOR(type,nr,size)   _IOC(_IOC_READ,(type),(nr),(_IOC_TYPECHECK(size)))
#define _IOW(type,nr,size)   _IOC(_IOC_WRITE,(type),(nr),(_IOC_TYPECHECK(size)))
#define _IOWR(type,nr,size)  _IOC(_IOC_READ|_IOC_WRITE,(type),(nr),(_IOC_TYPECHECK(size)))
```


I/O Control

This table lists ioctls visible from user land for Linux/x86. It contains most drivers up to 2.6.31, but I know I am missing some. There has been no attempt to list non-X86 architectures or ioctls from drivers/staging/.

Code	Seq#(hex)	Include File	Comments
=====			
0x00	00-1F	linux/fs.h	conflict!
0x00	00-1F	scsi/scsi_ioctl.h	conflict!
0x00	00-1F	linux/fb.h	conflict!
0x00	00-1F	linux/wavefront.h	conflict!
0x02	all	linux/fd.h	
0x03	all	linux/hdreg.h	
0x04	D2-DC	linux/umsdos_fs.h	Dead since 2.6.11, but don't reuse these.
0x06	all	linux/lp.h	
0x09	all	linux/raid/md_u.h	
0x10	00-0F	drivers/char/s390/vmcp.h	
0x10	10-1F	arch/s390/include/uapi/sclp_ctl.h	
0x10	20-2F	arch/s390/include/uapi/asm/hypfs.h	
0x12	all	linux/fs.h	
		linux/blkpg.h	
0x1b	all	InfiniBand Subsystem	< http://infiniband.sourceforge.net/ >
0x20	all	drivers/cdrom/cm206.h	
0x22	all	scsi/sg.h	
'#'	00-3F	IEEE 1394 Subsystem	Block for the entire subsystem

Agenda

I. Introduction

II. Device Drivers

I. I/O Control

II. Talking to a driver

III. Bug Hunting

I. Discovery & Debugging

II. Tooling

IV. Conclusion

Determining Access

- **/dev/random** is open to everyone
- **ppp** though, not so much

```
crw-rw-rw-    1 root root      1,   3 Jul 18 09:36 null
crw-----    1 root root      1,  12 Jul 18 09:36 oldmem
crw-r-----T  1 root kmem     1,   4 Jul 18 09:36 port
crw-----T   1 root root    108,   0 Jul 18 09:36 ppp
crw-----    1 root root     10,   1 Jul 18 09:36 psaux
crw-rw-rw-    1 root root      5,   2 Jul 24 14:22 ptmx
drwxr-xr-x    2 root root          0 Jul 18 09:36 pts
crw-rw-rw-    1 root root      1,   8 Jul 18 09:36 random
crw-rw-r-T+   1 root root    10,  59 Jul 18 09:36 rfkill
```

Capabilities

- Intended to split root into different privileges
 - No more all-or-nothing
- ~40 different caps can be applied to binaries
 - Eg. Want to sniff packets with Wireshark?
 - **setcap CAP_NET_RAW+ep /usr/bin/dumpcap**

Finding IOCTLs

- VirtualBox's SUPDrvIOC.h

```
/** Fast path IOCTL: VMMR0_DO_RAW_RUN */
#define SUP_IOCTL_FAST_DO_RAW_RUN          SUP_CTL_CODE_FAST(64)
/** Fast path IOCTL: VMMR0_DO_HM_RUN */
#define SUP_IOCTL_FAST_DO_HM_RUN          SUP_CTL_CODE_FAST(65)
/** Just a NOP call for profiling the latency of a fast ioctl call to VMMR0. */
#define SUP_IOCTL_FAST_DO_NOP            SUP_CTL_CODE_FAST(66)
```

```
#elif defined(RT_OS_LINUX)
    /* No automatic buffering, size limited to 16KB. */
    #include <linux/ioctl.h>
    #define SUP_CTL_CODE_SIZE(Function, Size)    _IOC(_IOC_READ | _IOC_WRITE, 'V', (Function) | SUP_IOCTL_FLAG, (Size))
    #define SUP_CTL_CODE_BIG(Function)          _IO('V', (Function) | SUP_IOCTL_FLAG)
    #define SUP_CTL_CODE_FAST(Function)         _IO('V', (Function) | SUP_IOCTL_FLAG)
    #define SUP_CTL_CODE_NO_SIZE(uIOctl)        ((uIOctl) & ~IOCSIZE_MASK)
```

Finding the IOCTL Handler

- VirtualBox's SUPDrv-linux.c

```
/** The file_operations structure. */
static struct file_operations gFileOpsVBoxDrvSys =
{
    owner:      THIS_MODULE,
    open:       VBoxDrvLinuxCreateSys,
    release:    VBoxDrvLinuxClose,
#ifdef HAVE_UNLOCKED_IOCTL
    unlocked_ioctl: VBoxDrvLinuxIOCtl,
#else
    ioctl:      VBoxDrvLinuxIOCtl,    <-- Oh, HAI!
#endif
};
```

ioctl() vs unlocked_ioctl()

- unlocked_ioctl() was added to kernel 2.6.11
 - IOCTLS no longer use the old BKL (Big Kernel **Lock**)
 - One less parameter (inode)
 - Not important for bug hunting, but now you know

```
int (*ioctl) (struct inode *inode, struct file *filp,  
              unsigned int cmd, unsigned long arg);
```

```
long (*unlocked_ioctl) (struct file *filp, unsigned int cmd,  
                        unsigned long arg);
```


compat_ioctl()

- Handles 32-bit processes calling ioctl() on 64-bit platforms
- Added in same kernel as unlocked_ioctl()

```
long (*compat_ioctl) (struct file *filp, unsigned int cmd,  
                     unsigned long arg);
```


Identify buffer types

- VirtualBox's SUPDrv-linux.c

```
static int VBoxDrvLinuxIOCtlSlow(struct file *pFilp, unsigned int uCmd, unsigned long ulArg, PSUPDRVSESSION pSession)
{
    int                rc;
    SUPREQHDR          Hdr;
    PSUPREQHDR          pHdr;
    uint32_t            cbBuf;

    Log6(("VBoxDrvLinuxIOCtl: pFilp=%p uCmd=%#x ulArg=%p pid=%d/%d\n", pFilp, uCmd, (void *)ulArg, RTProcSelf(), current->pid));

    /*
     * Read the header.
     */
    if (RT_UNLIKELY(copy_from_user(&Hdr, (void *)ulArg, sizeof(Hdr))))
    {
        Log(("VBoxDrvLinuxIOCtl: copy_from_user(,%#lx,) failed; uCmd=%#x.\n", ulArg, uCmd));
        return -EFAULT;
    }
}
```

Identify buffer types

- VBox's SUPDrvIOC.h

```
/**
 * Common In/Out header.
 */
typedef struct SUPREQHDR
{
    /** Cookie. */
    uint32_t      u32Cookie;
    /** Session cookie. */
    uint32_t      u32SessionCookie;
    /** The size of the input. */
    uint32_t      cbIn;
    /** The size of the output. */
    uint32_t      cbOut;
    /** Flags. See SUPREQHDR_FLAGS_* for details and values. */
    uint32_t      fFlags;
    /** The VBox status code of the operation, out direction only. */
    int32_t       rc;
} SUPREQHDR;
/** Pointer to a IOC header. */
typedef SUPREQHDR *PSUPREQHDR;
```



Kernel Debugging

- Host
 - Mac OS X (Yosemite)
 - VMware Fusion
- Guest
 - Ubuntu 14.04 (x86)

Kernel Debugging

- Configure the host
 - Copy the ISO, create a new VM, install guest OS
 - Edit the VMX config
 - (right-click and show package contents on Mac)
 - Add this line to the end of the file

```
debugStub.listen.guest32 = 1
```
 - Boot up guest VM

Kernel Debugging

- Configure the guest

```
codename=$(lsb_release -c | awk '{print $2}')
```

```
sudo tee /etc/apt/sources.list.d/ddebs.list << EOF
deb http://ddebs.ubuntu.com/ ${codename} main restricted univer
se multiverse
deb http://ddebs.ubuntu.com/ ${codename}-security main restrict
ed universe multiverse
deb http://ddebs.ubuntu.com/ ${codename}-updates main restricte
d universe multiverse
deb http://ddebs.ubuntu.com/ ${codename}-proposed main restrict
ed universe multiverse
EOF
```

```
sudo apt-key adv --keyserver keyserver.ubuntu.com --recv-keys
ECDCAD72428D7C01
```

```
sudo apt-get install linux-image-`uname -r`-dbg
```

Kernel Debugging

- Configure the guest
 - Copy the debug kernel to the host
 - `/usr/lib/debug/boot/vmlinux-<kernel version>-generic`

Kernel Debugging

- GDB

```
wget http://ftp.gnu.org/gnu/gdb/gdb-7.8.tar.gz
```

```
tar xf gdb-7.8.tar.gz
```

```
cd gdb-7.8
```

```
./configure --build=x86_64-apple-darwin14.0.0 --target=x86_64-vfs-linux --with-python && make
```

```
make install
```


Kernel Debugging

- GDB

```
(gdb) target remote :8832
```

```
Remote debugging using :8832
```

```
(gdb) symbol-file vmlinux-<kernel version>-generic
```

```
Reading symbols... done
```

Kernel Debugging

```
^C
Program received signal SIGINT, Interrupt.
0xc1147b1c in copy_page (from=<optimized out>, to=0xffffb9000)
    at /build/buildd/linux-3.13.0/arch/x86/include/asm/page_32.
h:47
47    /build/buildd/linux-3.13.0/arch/x86/include/asm/page_32.h
: No such file or directory.

(gdb) i r
eax          0xffffb9000    -290816
ecx          0x400        1024
edx          0x1000       4096
ebx          0x1e8        488
esp          0xedb4be48    0xedb4be48
ebp          0xedb4be68    0xedb4be68
esi          0xecbe8000    -323059712
edi          0xffffb9000    -290816
eip          0xc1147b1c    0xc1147b1c <copy_user_huge_page+76>
>
eflags       0x210246      [ PF ZF IF RF ID ]
cs           0x60         96
ss           0x68         104
ds           0x7b         123
es           0x7b         123
fs           0xd8         216
gs           0xe0         224

(gdb) c
Continuing.
```

Kernel Debugging

- Intentionally trigger a kernel panic
 - `echo c > /proc/sysrq-trigger`



Tooling

port·man·teau

/ˌpɔrtˈman(t)ō/

noun

a large trunk or suitcase, typically made of stiff leather and opening into two equal parts.

- consisting of or combining two or more separable aspects or qualities.



Tooling

- Experimental *nix tooling
 - Runs on Linux, but easily ported to other unix
- Highlights
 - Static code-based discovery and import of IOCTLs
 - SQLite database support
 - Generational fuzzing
 - Auto-generates PoCs

Tooling

```
[Portmanteau] v1.0
```

```
General>
```

```
[-D device]    - Set device path  
                Eg. -D /dev/net/tun
```

```
[-N n]         - Set number of iterations  
                Eg. -N 100000
```

```
[-z]           - Start a fuzzing run (typed-based random generation)
```

```
Database>
```

```
[-i file]      - Import IOCTL definitions from a single file  
                Eg. -i /path/to/driver_ioctl.h
```

```
[-a signature] - Add a new device signature to the database (manual)  
                Eg. -L -a "/dev/net/lsa:SNIFF_ENTIRE_INTERNET:0xdeadb33f:unsigned int"
```

```
[-d signature] - Delete a device signature from the database  
                Eg. -d "/dev/net/boring:IOCTL_EAT_CAKE"
```

Tooling

Driver Signatures

device_name	ioctl_name	ioctl_macro	ioctl_buftype
-----	-----	-----	-----
/dev/net/tun	TUNSETNOCSUM	_IOW(T,200,int)	unsignedshort
/dev/net/tun	TUNSETDEBUG	_IOW(T,201,int)	unsignedshort
/dev/net/tun	TUNSETIFF	_IOW(T,202,int)	unsignedshort
/dev/net/tun	TUNSETPERSIST	_IOW(T,203,int)	unsignedshort
/dev/net/tun	TUNSETOWNER	_IOW(T,204,int)	unsignedshort

Tooling

- Will make it available after the talk
 - Ping me if you'd like it sooner

Agenda

I. Introduction

II. Device drivers

I. I/O Control

II. Talking to a driver

III. Bug Hunting

I. Discovery & Debugging

II. Tooling

IV. Conclusion

Conclusion

- Drivers are a growing area for vuln research
 - Kernel code is becoming a more attractive target
 - Impact grows larger per app dependencies
- Fundamentals and tooling can help
- There's going to be bugs here for a long time
 - Think “forgotten syscalls”

Future Work

- Targeting abstraction layers
 - Eg. LibUSB
- Binary Analysis
 - IDA scripts that find & document IOCTL calls
- Other platforms
 - Diversity of bugs can vary based on how the kernel works

Thank you!

Questions?